

Executive Summary

Wearable physiological signals and behavioral cues can reliably indicate mental state ¹ ². For example, stress and anxiety typically raise heart rate and skin conductance ³ ⁴, while depressed or isolated individuals often show reduced variability (lower HRV) and blunted arousal ³ ⁵. Likewise, facial expressions (e.g. sadness, reduced movement) and social-media language patterns (negative affect, rumination) have been statistically linked to depression and loneliness ² ⁶. These findings motivate an integrated IoT+AI monitoring system: a wrist-band sensor records pulse, skin conductance and temperature, while a companion app analyzes live video (facial cues and movement) and optionally social-media activity. State-of-the-art machine-learning models (e.g. convolutional neural networks, LSTMs, transformers) can fuse these multimodal inputs to infer anxiety/depression/loneliness in near-real time ⁷ ⁸. We propose a low-latency architecture (shown below) that processes sensor and video data on-device and in the cloud, applies privacy-preserving analytics, and provides both patients and clinicians with interactive dashboards. The design balances accuracy (citing primary research) with safety, consent and compliance (GDPR/HIPAA), and includes a plan for clinical validation.

Clinical Evidence Linking Biosignals and Cues

- **Physiological signals:** Anxiety and acute stress trigger sympathetic overactivation, which manifests as *increased heart rate* and *increased galvanic skin response (GSR)* ³ ⁴. For example, an umbrella review notes that anxiety disorders show higher resting heart rate (sympathetic tone) ³, and laboratory studies find stress consistently elevates skin conductance ⁴. Conversely, chronic depression often involves *reduced parasympathetic activity* (lower HRV) ³. In one controlled study, isolated (low-arousal) subjects had *reduced HRV and lower GSR* compared to socializing or engaged states ⁵, suggesting these biosignals can index loneliness or social disengagement. (Notably, body/wrist temperature changes under stress are subtler, but a wearable thermometer can track thermoregulatory shifts or fever that may accompany mood changes.)
- **Facial expressions and movement:** Depressed individuals tend to exhibit characteristic facial cues. One study found “*sad expression and a smaller range of body movements*” in depressed patients ⁹. Machine-learning analysis of video has achieved high accuracy: a CNN+FACS system reported ~**87%** accuracy detecting depression and ~78% for anxiety from facial action units ². Similarly, decreased expressivity and micro-expressions can be detected by CNN/LSTM models ⁷. In practice, models integrating facial and head/gesture features yield strong correlation with clinical scales (e.g. >74% Pearson similarity to PHQ-9/GAD-7 scores ¹⁰).
- **Social media and digital behavior:** Language and usage patterns on Facebook/YouTube correlate with mental health. Depressed users often post fewer positive items and more negative/self-focused content; e.g. Facebook posts containing pain, rumination, personal pronouns predict higher depression scores ⁶. Loneliness manifests in social media with more cognitive-processing language (e.g. tentative, differentiation words) and fewer social-relationship mentions ¹¹. Heavy use of platforms is itself a marker: one study reported that “*frequent users of YouTube have higher levels of loneliness, anxiety, and depression*” ¹², and a review confirms *higher loneliness scores* among

heavy YouTube watchers ¹³. We will leverage such insights by analyzing a user's permitted social-media data (text posts, watch history keywords) to personalize the AI companion and cross-check predicted mood.

Sensor Specifications & Feasibility

- **Sensors:** A wrist-worn band can combine a photoplethysmogram (PPG) or ECG for pulse/HRV, a GSR module for skin conductance, and a skin temperature sensor. *Sampling rates:* Research-grade systems (e.g. Shimmer3) sample ECG at ~1024 Hz (24-bit) and PPG/GSR at ~128 Hz ¹⁴ ¹⁵. The Empatica E4 wristband, used in anxiety research, samples EDA and skin temperature at 4 Hz ¹⁶ ¹⁷. We recommend at least 50–100 Hz for PPG (for heart rate variability) and ~10–20 Hz for GSR (to capture rapid sweat events), although lower rates (4–10 Hz) may suffice for tonic-level monitoring ¹⁶. *Accuracy:* Typical PPG accuracy is ± 2 –5 beats per minute (cardiac). GSR sensors are calibrated to measure ~0.01–100 μ S conductivity ¹⁶. Skin temperature sensors (contact) have ~0.1–0.2°C accuracy ¹⁷.
- **Power & Connectivity:** A comfortable band can house a Li-ion battery supporting ~8–24-hour operation depending on sample rates and radio use (Bluetooth LE). Communication uses Bluetooth Low Energy (BLE) or IEEE 802.15.4 to the user's phone. For example, Phadke *et al.* built a GSR wristband with BLE (nRF52840 MCU) and demonstrated reliable streaming of stress data ¹⁸. Data is timestamped and sent in small packets (JSON or binary) over BLE; on the phone it can be buffered or relayed via WiFi/Cell to the cloud at configurable intervals.
- **Calibration:** Sensors require initial and periodic calibration. For GSR, dry electrodes need stable skin contact; filtering (e.g. Kalman filters) is used to normalize baseline drift ⁴. PPG requires adjustable LED drive for ambient light and motion artifacts; a few-second rest period can calibrate heart rate baseline. The system should allow the user to initiate a quick zeroing (e.g. handshake gesture calibration) if readings seem offset.

Signal Processing & AI Models

- **Preprocessing & Features:** Raw data will be denoised (e.g. band-pass filtering for ECG/PPG, smoothing for GSR, optical flow filtering for video). Key features include time-domain HRV metrics (RMSSD, SDNN), GSR phasic peaks (skin conductance responses), and statistical features (mean, variance). For video, facial landmarks and head pose are extracted (e.g. via dlib or MediaPipe), from which Action Units or emotion embeddings are computed. Gesture features (body motion, fidgeting rate) can be derived from frame differencing or pose estimation.
- **Model architectures:** We will employ a **multimodal deep learning** approach. For example, separate CNNs (e.g. ResNet or VGG variants) can process face images or optical flow streams to classify expressions ². Time-series signals (PPG, GSR) can feed into recurrent networks (LSTM/GRU) or 1D-CNNs. A late-fusion layer combines modalities, possibly via attention or transformer-based fusion, to output anxiety/depression levels. Established pretrained models (e.g. AffectNet-tuned CNNs for emotion, OpenFace for AUs, BERT/GPT for text sentiment) will be fine-tuned. For the AI chat, a large language model (GPT or similar) can be personalized on the user's social media text, with additional emotion conditioning.

- **Training data & metrics:** We will need labeled datasets pairing signals with clinical mood scores. Available corpora include DAIC-WOZ (videos with depression labels), AVEC (audio/video), and multimodal stress datasets. Transfer learning from general emotion datasets (AffectNet for faces, WESAD for physiological stress) will help. Evaluation metrics include classification accuracy/AUC and F1 for anxiety/depression episodes, and regression metrics (Pearson correlation, RMSE) when predicting severity scores ¹⁰. We will use cross-validation and hold-out test sets, and measure real-time inference latency. For example, Gavrilescu & Vizireanu achieved 87.2% accuracy on depression vs 77.9% on anxiety using a three-layer CNN/SVM network ², indicating feasible performance targets.

Privacy, Consent & Ethics

- **Data privacy:** All collected data is *sensitive health information*. Under GDPR, inferred mental health status is “special category” data requiring explicit consent ¹⁹. The system will require users to opt in to each data type (biosignals, video, social-media) and will allow them to revoke consent and delete data. Data at rest and in transit must be encrypted (e.g. TLS for network, AES-256 storage). Personally identifying information (name, social accounts) will be stored separately with strong pseudonymization. The system will adhere to HIPAA rules in the US for any clinical deployment: AWS, Azure, or GCP components must be HIPAA-eligible ²⁰ and we will sign required Business Associate Agreements.
- **Social media APIs & terms:** Integration uses only official APIs. For YouTube, we would use Google’s OAuth (scope `youtube.readonly`) to fetch watch history or liked videos. For Facebook, we use the Graph API with user authorization (permissions like `user_posts`, `user_likes`) and ensure compliance with Meta’s privacy policies. All social-media processing happens only on data the user explicitly permits; we store only derived, anonymized features (e.g. topic vectors), not raw content beyond retention needs. Ethical guidelines (e.g. informed consent, transparency) will be followed; for example, the app will inform users how AI companion responses are generated and will not make any clinical claims without a therapist’s oversight.
- **User control & transparency:** The app interface will include clear privacy settings (what data is collected and why). Clinicians accessing the data must also respect confidentiality. Potential risks (false alerts, bias) will be disclosed. Institutional review (IRB) will oversee the design, and the system will conform to frameworks like the Belmont Report and GDPR’s data minimization.

System Architecture & Data Flow

```
graph LR
    Wearable("Wearable Band\n(Temp, Pulse, GSR)") -->|BLE| PhoneApp("Mobile App")
    PhoneApp -->|Internet/WiFi| Cloud("Cloud Backend & ML")
    VideoCam("Camera (Video Stream)") --> PhoneApp
    SocialAPIs("Social Media API") --> Cloud
    Cloud -->|ML Inference| AIChat("AI Voice Companion")
    Cloud --> Clinician("Clinician Dashboard/Web Portal")
```

```
Cloud --> DB[(Secure Database)]  
PhoneApp --> DB
```

Wearable sensors transmit encrypted data via BLE to the mobile app, which forwards it (REST/HTTPS) to the cloud. Real-time facial video is analyzed on-device or streamed to edge compute. Social-media features are fetched via APIs and sent to the cloud. In the cloud, data enters a privacy-controlled database and triggers ML pipelines. The inference results feed two outputs: an AI chat companion that interacts with the user (voicebot) and a clinician portal. Communication uses HTTPS/TLS; data at rest is encrypted. To minimize latency, critical inference (facial emotion recognition) may be run on-device or at a nearby edge server, and model predictions are cached on the phone for offline review if connectivity drops. The architecture supports horizontal scaling (e.g. AWS IoT Core, Azure IoT Hub) and load-balancing ML inference (e.g. AWS SageMaker, GCP AI Platform).

UI/UX Considerations

The user app will have a **calm, engaging interface** to encourage regular use. Key features include:

- **Clear feedback:** A simple dashboard shows current mood scores (e.g. anxiety index), trend graphs, and daily log (with minimal text). Visualizations (color-coded gauges or charts) help users track progress.
- **Reminders & gamification:** Gentle prompts (e.g. “take a deep-breath break now”) and positive reinforcement (badges for consistency) can boost engagement. However, notifications will be infrequent to avoid annoyance.
- **Voice companion UI:** The AI chat interface should look friendly (e.g. a smiling avatar) and use empathetic language. It should let users skip or stop interaction at any time.
- **Accessibility:** The app will support large text, voice control, and color-blind-friendly palettes.
- **Clinician portal:** A web dashboard for therapists with aggregate data: longitudinal plots of physiological signals and inferred mood, flagged events (e.g. detected panic episodes), and exportable summaries. The design ensures clinicians can quickly see trends (weekly chart of GAD-7 predictions, for instance) and drill down (view raw signal around a stress spike). Role-based access control ensures only authorized clinicians see patient data.

Overall, usability will follow ISO 9241 guidelines (e.g. low cognitive load, intuitive navigation). We will prototype UI flows and test them with target users to refine clarity and reduce stigma.

Validation Plan (Clinical Study)

We propose a two-phase validation. **Phase 1 (Pilot):** Enroll ~30–50 participants with self-reported social anxiety or mild depression. Over 4–8 weeks, collect continuous data via the system. Use standardized scales (PHQ-9 for depression, GAD-7 for anxiety, UCLA Loneliness Scale) at baseline and weekly. Participants also keep brief mood diaries. We will evaluate: (a) **Convergent validity:** correlation between system scores (e.g. predicted anxiety index) and questionnaire results, targeting Pearson $r > 0.5$; (b) **Detection accuracy:** sensitivity/specificity of detecting high-anxiety episodes vs self-report logs. **Phase 2 (Controlled Trial):** A randomized trial ($N \approx 100$) comparing outcomes in an intervention group (receiving feedback and AI companion) vs control (minimal monitoring). Primary outcome: change in PHQ-9/GAD-7 at 8 weeks. Secondary: user engagement metrics, reduction in negative events (self-harm ideation, etc.). Sample size

will be powered (e.g. 80%) to detect moderate effect (Cohen's $d \sim 0.5$). Statistical analysis: mixed-effects models for repeated measures, chi-square tests for categorical outcomes, and survival analysis for time-to-event (e.g. panic attacks). Ethical approval will include informed consent detailing data use.

Deployment, Scalability & Maintenance

We envision a cloud-based deployment using a HIPAA-compliant provider. For instance:

- **AWS:** IoT Core for device messaging, Lambda or Greengrass for edge logic, SageMaker/EC2 for ML, DynamoDB/S3 for data (AWS is HIPAA-eligible ²⁰).
- **Azure:** IoT Hub, Azure Functions, Azure ML, Cosmos DB (HIPAA-ready services).
- **GCP:** Cloud IoT Core, Cloud Run/AI Platform, Firestore/BigQuery.

A multi-tenant, microservices architecture will allow scaling to many users. CI/CD pipelines (GitHub Actions, Terraform) will automate updates. Models will be retrained periodically as more data accrues. Monitoring tools (Prometheus, CloudWatch) will track system health and latency (aiming <200ms response for chat companion replies). We will provide an app-store-ready mobile client (iOS and Android) and a responsive web UI. Maintenance will include regular security audits (pen tests) and compliance reviews (ensuring new data policies are respected).

Risks, Limitations & Mitigation

- **Accuracy & bias:** Physiological signals vary across individuals (e.g. baseline heart rate by age or medication) and cultures (expressiveness differs). There is risk of false positives (misidentifying normal fluctuation as anxiety) or false negatives (masked stress). *Mitigation:* include personalization (user-specific baselines), conservative alert thresholds, and human-in-the-loop review (data seen by clinician before any action).
- **Privacy breaches:** Storing social-media or health data poses breach risk. *Mitigation:* end-to-end encryption, anonymization (store only feature vectors when possible), and intrusion detection. Data access logs and quick user-controlled delete functions are essential.
- **User adherence:** Users might stop wearing the device or skip video sessions, leading to missing data. *Mitigation:* design unobtrusive wearables (comfortable, aesthetically neutral), minimal required frequency (e.g. daily 5-min check-in), and backup modes (smartphone-only, with phone sensors as proxies).
- **Ethical misuse:** Automated mood inference could be stigmatizing or misused by third parties. *Mitigation:* strict access controls, only summarizing health metrics (not “your depression score is 8/10” to unauthorized viewers), and compliance with medical device regulations if needed (e.g. FDA guidance on mental health apps).
- **Regulatory:** Claiming diagnostic ability triggers regulation. *Mitigation:* position system as “decision support” or wellness tool, with clear disclaimers that it does not replace professional diagnosis.

Comparison of Alternatives

Component	Options	Pros	Cons
Pulse sensor	PPG (wrist)	Non-invasive, low-cost	Affected by motion/artifacts

Component	Options	Pros	Cons
GSR sensor	ECG (chest strap)	Gold-standard accuracy	Bulkier, less wearable
	PPG (ring)	Comfortable, constant contact	Fewer signals (no motion data)
	Finger electrodes (dry)	High sensitivity	Requires skin contact, sweat
	Wristband contacts	Wearable form-factor	Often lower signal amplitude
Temp sensor	Skin-contact thermistor	Simple, low-power	Influenced by environment
	Core (ear/forehead IR)	Reflects actual body temp	Harder to integrate
Facial AI	CNN (e.g. ResNet, MobileNet)	High accuracy (ImageNet pretraining) ²	Computationally heavy (edge challenges)
	Facial landmarks + heuristic	Lightweight	Less flexible, hand-tuned
Sequence Model	LSTM/GRU for time-series	Captures temporal patterns	Sensitive to long context
	1D-CNN for signals	Efficient, less memory	Fixed receptive field
Fusion	Early fusion (concatenate inputs)	Simpler	Modality imbalance
	Late fusion (combine outputs)	Modular	May lose cross-modal cues
	Attention/Transformer	State-of-art performance	Large data needed, complex
Cloud Provider	AWS	Mature IoT/ML stack, HIPAA-ready ²⁰	Complex pricing, vendor lock-in
	Azure	Good integration with MS services, compliance	Less global coverage in some regions
	GCP	Powerful AI tools (TPUs), BigQuery	Fewer IoT-specific features
	On-Premise/Edge	Data stays local (privacy)	Higher maintenance overhead

Each alternative involves trade-offs. For example, wrist PPG (as used by Empatica) sacrifices some accuracy for wearability ¹⁶ ; similarly, cloud vs edge decisions depend on latency needs and regulatory context.

Sources: We drew on peer-reviewed studies of biosignals and mental health ³ ⁵ ², sensor specifications from manufacturer documentation ¹⁴ ¹⁶, and regulatory guidance ¹⁹ ²⁰. The figures above and comparisons are synthesized from this literature.

¹ Using IoT architecture in health monitoring systems | Download Scientific Diagram

https://www.researchgate.net/figure/Using-IoT-architecture-in-health-monitoring-systems_fig2_371248558

² (PDF) Predicting Depression, Anxiety, and Stress Levels from Videos Using the Facial Action Coding System

[https://www.researchgate.net/publication/](https://www.researchgate.net/publication/335398761_Predicting_Depression_Anxiety_and_Stress_Levels_from_Videos_Using_the_Facial_Action_Coding_System)

[335398761_Predicting_Depression_Anxiety_and_Stress_Levels_from_Videos_Using_the_Facial_Action_Coding_System](https://www.researchgate.net/publication/335398761_Predicting_Depression_Anxiety_and_Stress_Levels_from_Videos_Using_the_Facial_Action_Coding_System)

³ Heart rate variability in mental disorders: an umbrella review of meta-analyses | Translational Psychiatry

https://www.nature.com/articles/s41398-025-03339-x?error=cookies_not_supported&code=bfb8caa3-6bf2-48e7-a36e-ab631a98c3d7

⁴ ¹⁸ Wearable IoT-Enabled Galvanic Skin Response Device for Objective Pain and Stress Monitoring: Hardware Design and Prototype Development

<https://www.mdpi.com/1424-8220/26/1/116>

⁵ (PDF) Smart Textile Based Wearable Sensors for Monitoring Isolation-Related Affective States: A Step Towards Loneliness Prediction

https://www.researchgate.net/publication/401615780_Smart_Textile_Based_Wearable_Sensors_for_Monitoring_Isolation-Related_Affective_States_A_Step_Towards_Loneliness_Prediction

⁶ A Systematic review of the validity of screening depression through Facebook, Twitter, Instagram, and Snapchat - ScienceDirect

<https://www.sciencedirect.com/science/article/abs/pii/S016503272100135X>

⁷ ⁸ Frontiers | Depression detection using deep learning and large language models from multimodalities

<https://www.frontiersin.org/journals/digital-health/articles/10.3389/fdgth.2026.1759857/full>

⁹ ¹⁰ Frontiers | Measuring depression severity based on facial expression and body movement using deep convolutional neural network

<https://www.frontiersin.org/journals/psychiatry/articles/10.3389/fpsyt.2022.1017064/full>

¹¹ Head versus heart: social media reveals differential language of loneliness from depression | npj Mental Health Research

https://www.nature.com/articles/s44184-022-00014-7?error=cookies_not_supported&code=efa725e3-f8a6-4b0f-8195-9419274ab8cd

¹² Impacts of YouTube on loneliness and mental health - Griffith News

<https://news.griffith.edu.au/2023/05/15/impacts-of-youtube-on-loneliness-and-mental-health/>

¹³ The Impact of YouTube on Loneliness and Mental Health

<https://www.mdpi.com/2227-9709/10/2/39>

¹⁴ ¹⁵ Sample Data - Shimmer Wearable Sensor Technology

<https://www.shimmersensing.com/support/sample-data/>

¹⁶ ¹⁷ Microsoft Word - E4_TechSpecs.docx

https://box.empatica.com/documentation/20141119_E4_TechSpecs.pdf

19 What is special category data? | ICO

<https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/special-category-data/what-is-special-category-data/>

20 HIPAA Compliance - Amazon Web Services (AWS)

<https://aws.amazon.com/compliance/hipaa-compliance/>